

Defense notes — Imagining the Target

Asma Hoseinpour Siouki · MSc Data Science · University of Padova · July 2026

**20-minute

1. Title

0:15 · cum 0:15

Good morning everyone, and thank you for being here.

2. Composed Image Retrieval

0:55 · cum 1:10

Imagine you see a chair online that you really like, and you want to find a dining set that includes that exact chair together with a table.

Searching with the image alone returns the same chair or similar chairs — but not as a dining set. Searching with text alone, "a dining table and chairs", returns many sets — but not your specific chair.

You need both: the image identifies the chair, the text describes the context you want to find it in.

This is composed image retrieval, the task of my thesis. The query is a reference image plus a modification text, and the goal is to retrieve images that preserve the reference object while satisfying the modification.

*Naming note: these are the two parts of the query, so from here on I call them the **query image** and the **query text**.*

3. Instance-level composed retrieval

0:40 · cum 1:50

My thesis focuses on the instance-level version of this task.

In category-level retrieval, any object of the correct class can be correct — any chair around a table would do. At the instance level, that is not enough: the system must retrieve the exact same chair shown in the query image, while also satisfying the query text.

So the challenge is to detect that one specific instance among many look-alikes, even when its context, viewpoint, or visual domain changes.

4. The i-CIR benchmark

0:50 · cum 2:40

We use the i-CIR benchmark, designed specifically for this task.

Each object instance has several composed queries — query image plus query text. Importantly, each query can have multiple correct targets: every database image showing the same object under the requested modification. So evaluation is about ranking all of them highly, not finding one right answer.

Each instance also comes with curated hard negatives: images that are visually similar but fail the modification, or satisfy the modification but show a different instance. This prevents the system from succeeding on one modality alone.

5. The i-CIR benchmark — scale

0:25 · cum 3:05

On scale: about 752 thousand database images, 202 instances, and 1,883 composed queries — a median of six per instance.

The histograms show the distributions, and confirm that many queries have more than one ground truth.

6. A zero-shot setting

0:40 · cum 3:45

We adopt a zero-shot setting — no training or fine-tuning on i-CIR.

Supervised CIR needs labelled triplets, and at the instance level every new object would require new annotation. A model trained on fixed instances also does not necessarily generalize to unseen ones — and new objects appear constantly. So supervision would hurt both scalability and generalizability.

With frozen pre-trained models, the system applies directly to unseen instances.

7. Problem formulation

0:20 · cum 4:05

Formally: given a query image and a query text, retrieve the database images relevant to both.

The system scores every database image and sorts them from most to least relevant. This ranking is the output — evaluation checks where the correct targets land in it.

8. Evaluation metric

0:35 · cum 4:40

Per query we use Average Precision, built on Precision-at-k — the fraction of the top-k results that are correct. AP is high when the correct images are near the top.

Averaging AP over all queries gives mean Average Precision.

But instances have different numbers of queries, so a direct average lets large instances dominate. Our main metric is therefore macro-mAP: average within each instance first, then across instances — every instance counts equally.

9. Vision-language models: CLIP and SigLIP

0:55 · cum 5:35

To represent images and text we use frozen vision-language models: CLIP and SigLIP.

Both are dual-encoders: an image encoder and a text encoder mapping into the same embedding space, where similarity is cosine similarity. We use them completely frozen — we only extract embeddings.

They differ in training. CLIP compares all image-text pairs within a batch through a softmax contrastive objective. SigLIP classifies each pair independently with a sigmoid loss — no global softmax, which scales better to large batches.

We compare both families to check that our method does not depend on one particular model.

10. The baseline: BASIC

1:00 · cum 6:35

Our baseline is BASIC, the zero-shot method introduced with the benchmark by Psomas and colleagues in 2025.

BASIC processes the two parts of the query separately: the query image goes through the visual encoder, the query text through the text encoder.

For every database image it computes two similarities — visual, to the query image, and textual, to the query text — and multiplies them. The multiplication acts like a soft logical

AND: a high score requires matching both parts.

BASIC uses CLIP ViT-L/14 as its main backbone — the standard in the zero-shot CIR literature — and we adopt the same one, so our results are directly comparable.

BASIC also includes further processing steps — centering, normalization, contextualization — which I will cover shortly.

11. Limitation of the baseline — our proposal

0:55 · cum 7:30

The key limitation: the query image and query text are never processed jointly. Each branch sees only half of the query, so the complete composed query is never explicitly represented.

Our proposal is a third branch. This is the full pipeline: as before, the image and text branches — and in addition, both query parts go together to a multimodal language model, which imagines the target and writes a caption of it. The caption is embedded by the same frozen text encoder and becomes a third similarity score.

The three scores are normalized and multiplied — a soft logical AND across the three branches, as the figure shows.

| *Fusion formulas and the Harris penalty → backup 24.*

12. The captioner: InternVL 3.5-8B

0:40 · cum 8:10

For caption generation we use InternVL 3.5, eight billion parameters, frozen — an open-source multimodal LLM that processes image and text jointly, at a practical balance of quality and cost. No fine-tuning anywhere.

In the chair example it sees the chair image and "placed around a table", and writes: "White metallic chair with scrolled armrests placed around a table."

Notice this single caption carries both the visual identity and the requested context — exactly the signal the two separate branches cannot represent.

13. Designing the instruction

0:35 · cum 8:45

Caption quality depends strongly on the instruction given to the model. I tested more than twenty instructions, varying the detail level, examples, constraint order, and restrictions against irrelevant information.

From these I selected five complementary strategies. Each describes the imagined target from a different perspective, and we average their caption embeddings.

| *What each strategy does → backup 32.*

14. Effect of the caption branch

0:45 · cum 9:30

A single caption improves macro-mAP by about 5.9 points, from 17.95 to 23.83 — it captures compositional information the two branches miss.

Averaging five captions from the five strategies goes further: it suppresses prompt-specific wording and keeps what is consistent. That reaches 25.41 — about 7.5 points over the image-text baseline.

| *Reminder on the slide: all results from here on are CLIP ViT-L/14 unless stated otherwise.*

15. Two refinements: centering and contextualisation 0:50 · cum 10:20

Two further refinements — one on the embeddings, one on the raw query text.

Centering: embedding spaces share a large common component — generic patterns rather than distinctive information. We subtract a per-modality mean computed on a separate corpus and re-normalise, isolating the semantic content.

Contextualisation acts on the raw query text — a query-side step, not score post-processing. The text encoder is trained on full captions, so a bare fragment like "sculpture" is out-of-distribution. We enrich the query with corpus terms before and after — "sculpture dog" — then embed and average the variants for a more robust text feature.

16. Adding the steps on top of each other 0:30 · cum 10:50

Adding each step cumulatively — baseline in blue, our caption pipeline in red, on CLIP-L.

The caption pipeline stays above the baseline along the entire ladder. Centering gives the baseline its big jump; contextualisation gives the largest single gain.

| *If asked why it stops here: projection and query expansion do not help once the caption branch is present → backup 25.*

17. Does it transfer? Four frozen backbones 0:45 · cum 11:35

We evaluate four frozen backbones: CLIP ViT-L — our main backbone — the larger CLIP ViT-H, SigLIP ViT-L, and the recent SigLIP2, the best-aligned of the four.

The caption branch improves every single one. The largest jump is SigLIP-L, where macro-mAP roughly doubles, from 21 to 43. SigLIP2 is the strongest base model.

So the contribution does not depend on one particular vision-language model.

18. Our best overall system 0:40 · cum 12:15

The final system: SigLIP2, the caption branch with five averaged captions, and centering plus contextualisation. Everything frozen.

It reaches 61.98 macro-mAP. Beyond mAP: the correct target is ranked first for 71% of queries and appears in the top ten for 95%. Recall at one hundred is 94%.

That high recall matters — I will come back to it in the future work.

| *The SigLIP2 ladder behind 61.98 → backup 26.*

19. Qualitative example — success 0:25 · cum 12:40 (flex)

A success: Homer Simpson, "as a plushie".

The plain product gets 4% AP — generic plushies: SpongeBob, minions. BASIC reaches 33%. The caption — "Homer Simpson as a plushie, same facial features" — pins the retrieval to actual Homer plushies: 88%. The caption keeps the identity while honouring the modification.

20. Qualitative example — failure 0:25 · cum 13:05 (flex)

An honest failure: gold metal eyeglasses, "worn by a dog".

All three methods score near zero. The single ground truth is a dog actually wearing these glasses; every method returns glasses on tables instead. When the modification demands a drastic, rarely-photographed context, the image branch dominates and the target never gets near the top.

21. Limitations and future work

0:40 · cum 13:45

Limitations: caption generation costs about fifteen seconds per query and dominates inference; results depend on prompt design — averaging mitigates but does not remove it; and the fixed product cannot re-weight the branches.

Future work: a more efficient captioner would cut the main cost. A learned fusion of the three scores could trade the zero-shot constraint for accuracy. And since recall is already high, a second-stage reranker is promising — most remaining error is in the ordering of candidates already retrieved, and a caption-conditioned reranker keeps the modification in the loop. Finally, the caption-as-third-modality idea is benchmark-agnostic and could be tested beyond i-CIR.

22. Conclusions

0:30 · cum 14:15

Two conclusions.

The caption is the single most effective component — it improves every backbone on the full dataset.

And SigLIP2 is the strongest backbone for the pipeline, reaching the thesis-best 61.98 macro-mAP, fully zero-shot.

If pushed: the caption alone beats the plain product on CLIP-L → backup 27; centering and the caption stack → backup 28.

23. Thank you

Thank you. I am happy to take questions.

Q&A map

Score fusion formulas + Harris penalty → **backup 24**. Full CLIP-L ablation, every rung (incl. projection and query expansion) → **backup 25**. SigLIP2 ladder (produces 61.98) → **backup 26**. Component contributions — image-only, text-only, caption-only, product, triplet → **backup 27**. Centering + caption are complementary → **backup 28**. "The paper reports 34.35 for BASIC" → **backup 29**. Latency breakdown → **backup 30**. Five instruction strategies, sample captions → **backup 31**; what each does → **backup 32**.

Problem formulation on slide 7; metric on slide 8; full pipeline figure on slide 11; backbone table on slide 17.

Backup-slide notes

24. Combining the three scores

Min-based normalization puts the three branches on comparable scales; negative residuals are clamped to zero — no positive evidence. The normalized scores are multiplied: a soft logical AND. On top, a Harris-inspired penalty, from the Harris corner detector: a candidate must not win on one branch alone — the exact chair not around a table, or tables with the wrong chair, are both suppressed. The penalty is weighted by a hyperparameter λ , fixed across all experiments. The surface shows the score is high only when both arguments are simultaneously large.

25. Full CLIP-L ablation

Every rung, for the three pipelines. Two extra BASIC steps appear here: projection and query expansion. Projection barely helps BASIC (32.25 → 32.48) and *hurts* the caption pipelines; query expansion hurts everything. That is why our pipeline stops at contextualisation. Best triplet: 35.76.

26. SigLIP2 ladder

Same shape as CLIP-L: the caption pipeline stays above BASIC throughout, and contextualisation gives the final jump. Best configuration 61.98 macro-mAP vs 57.69 for BASIC at the same rung — and 38.1 for the plain product.

27. Component contributions (CLIP-L)

Each branch alone, then the combinations — raw similarity, no post-processing. Image only 2.75, text only 3.26: neither works alone. The caption alone reaches 19.51 — it already beats the full image × text product at 17.95, because it is the only signal that sees both query parts. Adding it to the product: 25.41.

28. Centering and the caption are complementary

The caption gain is not absorbed by centering. At the raw product the caption adds +7.5; after centering it still adds +4.1. The two contributions stack, which is why the best configuration keeps both.

29. Reproducing BASIC

The paper reports 34.35; our re-run of the released code gives 32.48. The gap is concentrated in the image branch and is consistent with re-extracting features in a newer software stack — image decoding and resizing differ across library versions. Contextualisation also draws fresh object nouns per query. Both BASIC and our method run in the same environment, so every comparison in the thesis is unaffected.

30. Inference latency

100 queries, RTX-6000-Ada. Caption generation — five captions — takes about 14.7 seconds and is 99.7% of the online cost; embedding the query triplet is milliseconds, and the backbone choice does not change latency. Database embeddings are precomputed offline. Levers: fewer captions, shorter generations, a smaller or quantised MLLM.

31. The five instructions — sample captions

The five strategies on the Tintin example, "as an iconic rooftop sign". Each phrases the same target differently; their embeddings are averaged (avg-5).

32. The five instructions — explained

Adaptive precision adjusts the detail level to the query. Constraint-priority states the modification first so it is never dropped. The anti-noise contract forbids inventing attributes not visible in the image. The slot-fill template keeps captions uniform and comparable. Draft-refine adds a self-correction pass. They were chosen from twenty-plus candidates because each controls a different failure mode.